

## Homework 2

July 17, 2018

### 1 Chapter 4, Exercise 4 (p. 168)

#### 1.1 a

$$(0.95 - 0.05) * 0.1 + 2 \int_0^{0.05} (x + 0.05) dx = 0.09 + 0.0075 = 0.0975$$

(If we ignore  $X < 0.05$  and  $X > 0.95$ , the answer should be 0.1)

#### 1.2 b

Ignoring corner cases, the answer is 0.01.

#### 1.3 c

Ignoring corner cases, the answer is  $0.1^{100} = 10^{-100}$

#### 1.4 d

When  $p$  is large (compared to  $n$ ),  $n$  observations would be very sparse in the  $p$ -dimension space, making KNN unreliable.

#### 1.5 e

Suppose the length is  $l$ , then

$$l^p = 0.1$$

$$l = 0.1^{\frac{1}{p}}$$

- $p = 1, l = 0.100$
- $p = 2, l = 0.316$
- $p = 100, l = 0.977$

As  $p$  grows larger,  $l$  converges to 1, which means almost the whole space is needed to make a reasonable prediction.

## 2 Chapter 4, Exercise 6 (p. 170)

### 2.1 a

$$P(Y = 1|X_1, X_2) = \frac{e\beta_0 + \beta_1 X_1 + \beta_2 X_2}{1 + e\beta_0 + \beta_1 X_1 + \beta_2 X_2} = 0.3775$$

```
In [1]: x1 = 40
        x2 = 3.5
        b0 = -6
        b1 = 0.05
        b2 = 1
        p = exp(b0+b1*x1+b2*x2)/(1+exp(b0+b1*x1+b2*x2))
        p

0.377540668798145
```

### 2.2 b

Solve the equation about  $X_1$ :

$$P(Y = 1|X_1, X_2) = \frac{e\beta_0 + \beta_1 X_1 + \beta_2 X_2}{1 + e\beta_0 + \beta_1 X_1 + \beta_2 X_2} = 0.5$$

That is,

$$\log\left[\frac{P(Y = 1|X_1, X_2)}{P(Y = 0|X_1, X_2)}\right] = \beta_0 + \beta_1 X_1 + \beta_2 X_2 = 0$$

We have  $X_1 = 50$ .

```
In [2]: x2 = 3.5
        b0 = -6
        b1 = 0.05
        b2 = 1
        x1 = (-b2*x2-b0)/b1
        x1

50
```

## 3 Chapter 4, Exercise 8 (p. 170)

Performing 1-nearest neighbors on training set always produces correct result. Therefore the test error rate of 1-nearest neighbors =  $18\%/0.5 = 36\%$ . The error rate of logistic regression is 36%, so it is better for this case.

## 4 Chapter 4, Exercise 10 (p. 171)

```
In [3]: library(ISLR)
        data(Weekly)
```

Warning message:  
package ISLR was built under R version 3.4.2

## 4.1 a

```
In [4]: summary(Weekly)
        pairs(Weekly)
        round(cor(Weekly[, -9]), 3)
```

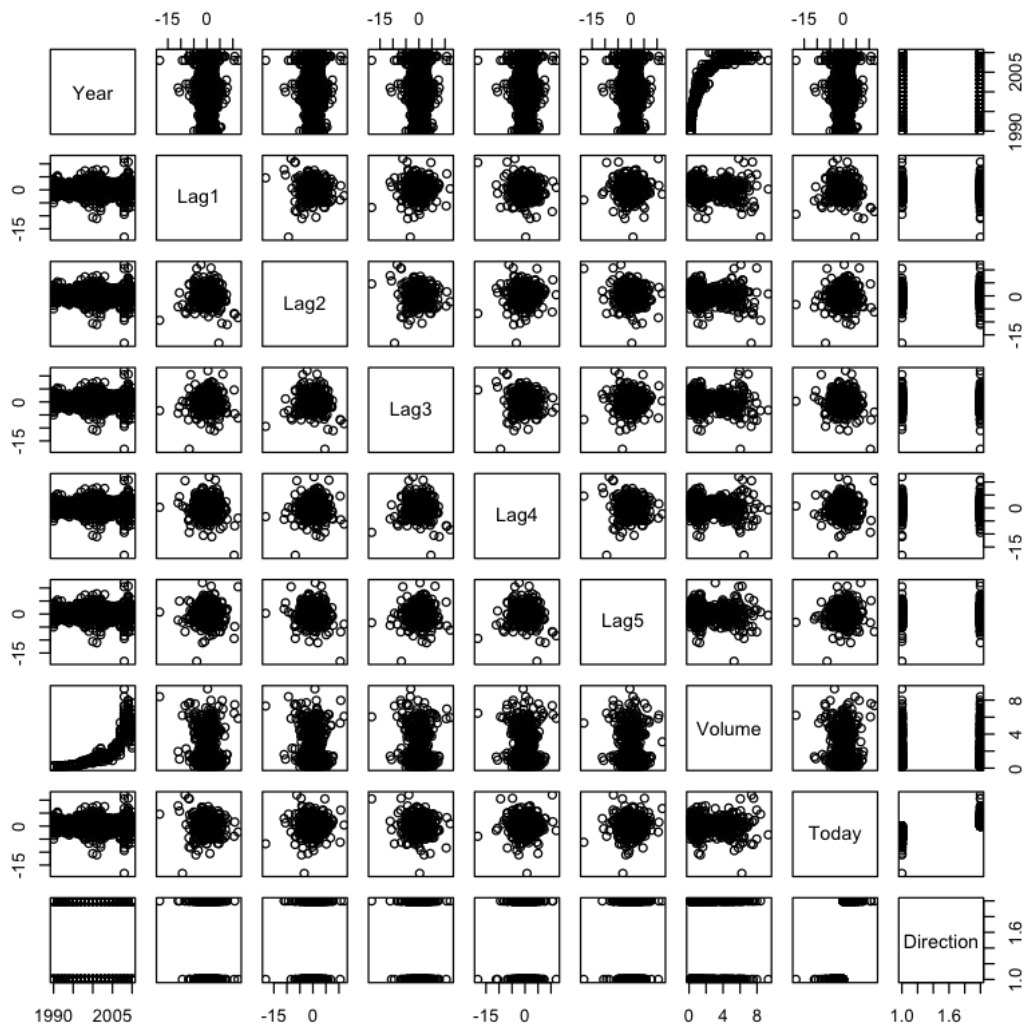
Year	Lag1	Lag2	Lag3
Min. :1990	Min. : -18.1950	Min. : -18.1950	Min. : -18.1950
1st Qu.:1995	1st Qu.: -1.1540	1st Qu.: -1.1540	1st Qu.: -1.1580
Median :2000	Median : 0.2410	Median : 0.2410	Median : 0.2410
Mean :2000	Mean : 0.1506	Mean : 0.1511	Mean : 0.1472
3rd Qu.:2005	3rd Qu.: 1.4050	3rd Qu.: 1.4090	3rd Qu.: 1.4090
Max. :2010	Max. : 12.0260	Max. : 12.0260	Max. : 12.0260

Lag4	Lag5	Volume	Today
Min. : -18.1950	Min. : -18.1950	Min. : 0.08747	Min. : -18.1950
1st Qu.: -1.1580	1st Qu.: -1.1660	1st Qu.: 0.33202	1st Qu.: -1.1540
Median : 0.2380	Median : 0.2340	Median : 1.00268	Median : 0.2410
Mean : 0.1458	Mean : 0.1399	Mean : 1.57462	Mean : 0.1499
3rd Qu.: 1.4090	3rd Qu.: 1.4050	3rd Qu.: 2.05373	3rd Qu.: 1.4050
Max. : 12.0260	Max. : 12.0260	Max. : 9.32821	Max. : 12.0260

Direction  
Down:484  
Up :605

	Year	Lag1	Lag2	Lag3	Lag4	Lag5	Volume	Today
Year	1.000	-0.032	-0.033	-0.030	-0.031	-0.031	0.842	-0.032
Lag1	-0.032	1.000	-0.075	0.059	-0.071	-0.008	-0.065	-0.075
Lag2	-0.033	-0.075	1.000	-0.076	0.058	-0.072	-0.086	0.059
Lag3	-0.030	0.059	-0.076	1.000	-0.075	0.061	-0.069	-0.071
Lag4	-0.031	-0.071	0.058	-0.075	1.000	-0.076	-0.061	-0.008
Lag5	-0.031	-0.008	-0.072	0.061	-0.076	1.000	-0.059	0.011
Volume	0.842	-0.065	-0.086	-0.069	-0.061	-0.059	1.000	-0.033
Today	-0.032	-0.075	0.059	-0.071	-0.008	0.011	-0.033	1.000



Year and Volume could have a positive correlation.

## 4.2 b

```
In [5]: fit1 = glm(Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 + Volume, data = Weekly,
family = binomial)
summary(fit1)
```

Call:

```
glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +
Volume, family = binomial, data = Weekly)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.6949	-1.2565	0.9913	1.0849	1.4579

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	0.26686	0.08593	3.106	0.0019 **
Lag1	-0.04127	0.02641	-1.563	0.1181
Lag2	0.05844	0.02686	2.175	0.0296 *
Lag3	-0.01606	0.02666	-0.602	0.5469
Lag4	-0.02779	0.02646	-1.050	0.2937
Lag5	-0.01447	0.02638	-0.549	0.5833
Volume	-0.02274	0.03690	-0.616	0.5377

---

Signif. codes: 0 \*\*\* 0.001 \*\* 0.01 \* 0.05 . 0.1 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1496.2 on 1088 degrees of freedom  
 Residual deviance: 1486.4 on 1082 degrees of freedom  
 AIC: 1500.4

Number of Fisher Scoring iterations: 4

Lag2 has more statistical significance.

### 4.3 c

```
In [6]: attach(Weekly)
        threshold = 0.5
        fit1.ans = predict(fit1, type = "response")
        fit1.preds = ifelse(fit1.ans > 0.5, "Up", "Down")
        table(fit1.preds, Direction)
        mean(fit1.preds == Direction)
```

	Direction
fit1.preds	Down Up
Down	54 48
Up	430 557

0.561065197428834

Overall fraction of correct predictions = 56.1%.

Most incorrect predictions mistake Down for Up ( $\frac{430}{430+48} = 90.0\%$ ).

```
In [7]: train = (Year <= 2008)
        train_data = Weekly[train,]
        test_data = Weekly[!train,]
```

```

fit2 = glm(Direction ~ Lag2, data = train_data, family = binomial)
fit2.ans = predict(fit2, test_data, type = "response")
fit2.preds = ifelse(fit2.ans > 0.5, "Up", "Down")
table(fit2.preds, test_data$Direction)
mean(fit2.preds == test_data$Direction)
detach(Weekly)

```

```

fit2.preds Down Up
      Down    9  5
      Up    34 56

```

0.625

Overall fraction of correct predictions for the held out data = 62.5%.

## 5 Chapter 5, Exercise 5 (p. 198)

```

In [8]: library(ISLR)
        attach(Default)
        # summary(Default)

```

### 5.1 a

```

In [9]: fit1 = glm(default ~ income + balance, data = Default, family = binomial)
        summary(fit1)

```

Call:

```

glm(formula = default ~ income + balance, family = binomial,
     data = Default)

```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.4725	-0.1444	-0.0574	-0.0211	3.7245

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-1.154e+01	4.348e-01	-26.545	< 2e-16 ***
income	2.081e-05	4.985e-06	4.174	2.99e-05 ***
balance	5.647e-03	2.274e-04	24.836	< 2e-16 ***

---

Signif. codes: 0 \*\*\* 0.001 \*\* 0.01 \* 0.05 . 0.1 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 2920.6 on 9999 degrees of freedom  
Residual deviance: 1579.0 on 9997 degrees of freedom

AIC: 1585

Number of Fisher Scoring iterations: 8

## 5.2 b

```
In [10]: TrainAndEvaluate = function(train_ratio) {  
  set.seed(2)  
  # i.  
  train = sample(nrow(Default), train_ratio*nrow(Default))  
  train_data = Default[train,]  
  validate_data = Default[-train,]  
  # ii.  
  fit = glm(default ~ income + balance, data = train_data, family = binomial)  
  # iii.  
  fit.ans = predict(fit, validate_data, type = "response")  
  fit.preds = ifelse(fit.ans > 0.5, "Yes", "No")  
  # iv.  
  return(mean(fit.preds != validate_data$default))  
}
```

TrainAndEvaluate(0.8)

0.0215

## 5.3 c

```
In [11]: TrainAndEvaluate(0.5)  
TrainAndEvaluate(0.6)  
TrainAndEvaluate(0.7)
```

0.0276

0.02725

0.024

As the ratio of training data increases, the validation error generally reduces.

## 5.4 d

```
In [12]: TrainAndEvaluate = function(train_ratio) {  
  set.seed(2)  
  # i.  
  train = sample(nrow(Default), train_ratio*nrow(Default))  
  train_data = Default[train,]  
  validate_data = Default[-train,]  
  # ii.  
  fit = glm(default ~ income + student + balance, data = train_data,  
            family = binomial)
```

```

    # iii.
    fit.ans = predict(fit, validate_data, type = "response")
    fit.preds = ifelse(fit.ans > 0.5, "Yes", "No")
    # iv.
    return(mean(fit.preds != validate_data$default))
}

```

```

TrainAndEvaluate(0.5)
TrainAndEvaluate(0.6)
TrainAndEvaluate(0.7)
TrainAndEvaluate(0.8)

```

```

0.0286
0.02825
0.025
0.022

```

Including a dummy variable for student does not lead to a reduction in the test error rate.

## 6 Chapter 5, Exercise 6 (p. 199)

### 6.1 a

```

In [13]: fit1 = glm(default ~ income + balance, data = Default, family = binomial)
          summary(fit1)

```

Call:

```

glm(formula = default ~ income + balance, family = binomial,
    data = Default)

```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.4725	-0.1444	-0.0574	-0.0211	3.7245

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-1.154e+01	4.348e-01	-26.545	< 2e-16 ***
income	2.081e-05	4.985e-06	4.174	2.99e-05 ***
balance	5.647e-03	2.274e-04	24.836	< 2e-16 ***

---

Signif. codes: 0 \*\*\* 0.001 \*\* 0.01 \* 0.05 . 0.1 1

(Dispersion parameter for binomial family taken to be 1)

```

Null deviance: 2920.6 on 9999 degrees of freedom
Residual deviance: 1579.0 on 9997 degrees of freedom
AIC: 1585

```



Number of Fisher Scoring iterations: 8

Estimated standard error:  $4.985e - 06$  for income,  $2.274e - 04$  for balance

## 6.2 b

```
In [14]: boot.fn = function(data, index) {  
  return(coef(glm(default ~ income + balance,  
    data = data, family = binomial, subset = index)))  
}
```

## 6.3 c

```
In [15]: library(boot)  
  set.seed(1)  
  boot(Default, boot.fn, 50)
```

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:

```
boot(data = Default, statistic = boot.fn, R = 50)
```

Bootstrap Statistics :

	original	bias	std. error
t1*	-1.154047e+01	1.181200e-01	4.202402e-01
t2*	2.080898e-05	-5.466926e-08	4.542214e-06
t3*	5.647103e-03	-6.974834e-05	2.282819e-04

## 6.4 d

There are no significant difference between the estimated standard errors produced by two methods.

## 7 Chapter 5, Exercise 8 (p. 200)


### 7.1 a

```
In [16]: set.seed(1)  
  y=rnorm(100)  
  x=rnorm(100)  
  y=x-2*x^2+ rnorm(100)
```

- $n = 100$  (observations)

- $p = 2$  (predictors)
- $Y = -2X^2 + X + \epsilon$  (model)

In [17]: `plot(x,y)`



Quadratic relationship between  $X$  and  $Y$  with noises.

## 7.2 c

```
In [18]: polyfits = function() {
  Data = data.frame(x, y)
  # i.
  glm.fit1 = glm(y ~ x)
  print(cv.glm(Data, glm.fit1)$delta)
```

```

    # ii.
    glm.fit2 = glm(y ~ poly(x, 2, raw=TRUE))
    print(cv.glm(Data, glm.fit2)$delta)
    # iii.
    glm.fit3 = glm(y ~ poly(x, 3, raw=TRUE))
    print(cv.glm(Data, glm.fit3)$delta)
    # iv.
    glm.fit4 = glm(y ~ poly(x, 4, raw=TRUE))
    print(cv.glm(Data, glm.fit4)$delta)
}

set.seed(1)
polyfits()

[1] 5.890979 5.888812
[1] 1.086596 1.086326
[1] 1.102585 1.102227
[1] 1.114772 1.114334

```

```

In [19]: set.seed(233)
         polyfits()

```

```

[1] 5.890979 5.888812
[1] 1.086596 1.086326
[1] 1.102585 1.102227
[1] 1.114772 1.114334

```

Results are same, because LOOCV predicts  $n$  observations separately using rest of the data, therefore no randomness is introduced.

### 7.3 e

The quadratic model has the smallest error, as we expected. The data are generated using quadratic model, so the model should perform better than others.

### 7.4 f

```

In [20]: glm.fit1 = glm(y ~ x)
         glm.fit2 = glm(y ~ poly(x, 2, raw=TRUE))
         glm.fit3 = glm(y ~ poly(x, 3, raw=TRUE))
         glm.fit4 = glm(y ~ poly(x, 4, raw=TRUE))
         summary(glm.fit1)
         summary(glm.fit2)
         summary(glm.fit3)
         summary(glm.fit4)

```

Call:

```
glm(formula = y ~ x)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-7.3469	-0.9275	0.8028	1.5608	4.3974

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-1.8185	0.2364	-7.692	1.14e-11 ***
x	0.2430	0.2479	0.981	0.329

---

Signif. codes: 0 \*\*\* 0.001 \*\* 0.01 \* 0.05 . 0.1 1

(Dispersion parameter for gaussian family taken to be 5.580018)

Null deviance: 552.21 on 99 degrees of freedom  
Residual deviance: 546.84 on 98 degrees of freedom  
AIC: 459.69

Number of Fisher Scoring iterations: 2

Call:

```
glm(formula = y ~ poly(x, 2, raw = TRUE))
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.89884	-0.53765	0.04135	0.61490	2.73607

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-0.09544	0.13345	-0.715	0.476
poly(x, 2, raw = TRUE)1	0.89961	0.11300	7.961	3.24e-12 ***
poly(x, 2, raw = TRUE)2	-1.86665	0.09151	-20.399	< 2e-16 ***

---

Signif. codes: 0 \*\*\* 0.001 \*\* 0.01 \* 0.05 . 0.1 1

(Dispersion parameter for gaussian family taken to be 1.06575)

Null deviance: 552.21 on 99 degrees of freedom  
Residual deviance: 103.38 on 97 degrees of freedom  
AIC: 295.11

Number of Fisher Scoring iterations: 2

Call:

```
glm(formula = y ~ poly(x, 3, raw = TRUE))
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.87250	-0.53881	0.02862	0.59383	2.74350

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-0.09865	0.13453	-0.733	0.465
poly(x, 3, raw = TRUE)1	0.95551	0.22150	4.314	3.9e-05 ***
poly(x, 3, raw = TRUE)2	-1.85303	0.10296	-17.998	< 2e-16 ***
poly(x, 3, raw = TRUE)3	-0.02479	0.08435	-0.294	0.769

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for gaussian family taken to be 1.075883)

Null deviance: 552.21 on 99 degrees of freedom  
Residual deviance: 103.28 on 96 degrees of freedom  
AIC: 297.02

Number of Fisher Scoring iterations: 2

Call:

```
glm(formula = y ~ poly(x, 4, raw = TRUE))
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.8914	-0.5244	0.0749	0.5932	2.7796

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-0.13897	0.15973	-0.870	0.386455
poly(x, 4, raw = TRUE)1	0.90980	0.24249	3.752	0.000302 ***
poly(x, 4, raw = TRUE)2	-1.72802	0.28379	-6.089	2.4e-08 ***
poly(x, 4, raw = TRUE)3	0.00715	0.10832	0.066	0.947510
poly(x, 4, raw = TRUE)4	-0.03807	0.08049	-0.473	0.637291

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for gaussian family taken to be 1.084654)

Null deviance: 552.21 on 99 degrees of freedom  
Residual deviance: 103.04 on 95 degrees of freedom  
AIC: 298.78

Number of Fisher Scoring iterations: 2

Linear and quadratic terms have statistical significance, as shown by the p-values. The summary results agree with the CV results.

## 8 Chapter 5, Exercise 9

```
In [21]: library(MASS)
         summary(Boston)
         attach(Boston)
```

crim	zn	indus	chas
Min. : 0.00632	Min. : 0.00	Min. : 0.46	Min. : 0.00000
1st Qu.: 0.08204	1st Qu.: 0.00	1st Qu.: 5.19	1st Qu.: 0.00000
Median : 0.25651	Median : 0.00	Median : 9.69	Median : 0.00000
Mean : 3.61352	Mean : 11.36	Mean : 11.14	Mean : 0.06917
3rd Qu.: 3.67708	3rd Qu.: 12.50	3rd Qu.: 18.10	3rd Qu.: 0.00000
Max. : 88.97620	Max. : 100.00	Max. : 27.74	Max. : 1.00000

nox	rm	age	dis
Min. : 0.3850	Min. : 3.561	Min. : 2.90	Min. : 1.130
1st Qu.: 0.4490	1st Qu.: 5.886	1st Qu.: 45.02	1st Qu.: 2.100
Median : 0.5380	Median : 6.208	Median : 77.50	Median : 3.207
Mean : 0.5547	Mean : 6.285	Mean : 68.57	Mean : 3.795
3rd Qu.: 0.6240	3rd Qu.: 6.623	3rd Qu.: 94.08	3rd Qu.: 5.188
Max. : 0.8710	Max. : 8.780	Max. : 100.00	Max. : 12.127

rad	tax	ptratio	black
Min. : 1.000	Min. : 187.0	Min. : 12.60	Min. : 0.32
1st Qu.: 4.000	1st Qu.: 279.0	1st Qu.: 17.40	1st Qu.: 375.38
Median : 5.000	Median : 330.0	Median : 19.05	Median : 391.44
Mean : 9.549	Mean : 408.2	Mean : 18.46	Mean : 356.67
3rd Qu.: 24.000	3rd Qu.: 666.0	3rd Qu.: 20.20	3rd Qu.: 396.23
Max. : 24.000	Max. : 711.0	Max. : 22.00	Max. : 396.90

lstat	medv
Min. : 1.73	Min. : 5.00
1st Qu.: 6.95	1st Qu.: 17.02
Median : 11.36	Median : 21.20
Mean : 12.65	Mean : 22.53
3rd Qu.: 16.95	3rd Qu.: 25.00
Max. : 37.97	Max. : 50.00

### 8.1 a

```
In [22]: medv.mean = mean(medv)
         medv.mean
```

22.5328063241107

### 8.2 b

```
In [23]: medv.err = sd(medv)/sqrt(nrow(Boston))
         medv.err
```

0.408861147497535

### 8.3 c

```
In [24]: set.seed(1)
         library(boot)
         boot.fn = function(data, index){
           return(mean(data[index]))
         }
         medv.bstrap = boot(medv, boot.fn, 1000)
         medv.bstrap
```

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:

```
boot(data = medv, statistic = boot.fn, R = 1000)
```

Bootstrap Statistics :

	original	bias	std. error
t1*	22.53281	0.008517589	0.4119374

The two results are very close (~0.7% difference).

### 8.4 d

```
In [25]: c(medv.bstrap$t0 - sd(medv.bstrap$t) * 2, medv.bstrap$t0 + sd(medv.bstrap$t) * 2)
         t.test(Boston$medv)
```

1. 21.708931444342 2. 23.3566812038793

One Sample t-test

data: Boston\$medv

```

t = 55.111, df = 505, p-value < 2.2e-16
alternative hypothesis: true mean is not equal to 0
95 percent confidence interval:
 21.72953 23.33608
sample estimates:
mean of x
 22.53281

```

The two results are very close (~0.08% difference).

## 8.5 e

```

In [26]: medv.med = median(medv)
          medv.med

21.2

```

## 8.6 f

```

In [27]: set.seed(1)
          library(boot)
          boot.fn.med = function(data, index){
            return(median(data[index]))
          }
          med.bstrap = boot(medv, boot.fn.med, 1000)
          med.bstrap

```

## ORDINARY NONPARAMETRIC BOOTSTRAP

Call:

```
boot(data = medv, statistic = boot.fn.med, R = 1000)
```

Bootstrap Statistics :

	original	bias	std. error
t1*	21.2	-0.01615	0.3801002

We get a result with reasonably small standard error.

## 8.7 g

```

In [28]: medv.tp = quantile(medv, 0.1)
          medv.tp

10\%: 12.75

```



## 8.8 h

```
In [29]: set.seed(1)
         library(boot)
         boot.fn.tentp = function(data, index){
           return(quantile(data[index], 0.1))
         }
         tentp.bstrap = boot(medv, boot.fn.tentp, 1000)
         tentp.bstrap
```

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:

```
boot(data = medv, statistic = boot.fn.tentp, R = 1000)
```

Bootstrap Statistics :

	original	bias	std. error
t1*	12.75	0.01005	0.505056

The standard error of estimating tenth percentile is larger then that of estimating median.